

A Tiny Taste of Machine Learning

All ML Methods Require

- Representation of the features
- Distance metric for feature vectors
- Objective function and constraints
- Optimization method for learning the model
- Evaluation method

Distance Between Vectors


- Minkowski metric

p-norm

$$\text{dist}(X1, X2, p) = \left(\sum_{k=1}^{\text{len}} \text{abs}(X1_k - X2_k)^p \right)^{1/p}$$

p = 1: Manhattan Distance 1-norm


P = 2: Euclidean Distance 2-norm



Manhattan (1902)



Which one is nearer?



Euclid-d: $\sqrt{8}$

Manhattan-d: 3

An Example

Features						Label
Name	Egg-laying	Scales	Poisonous	Cold-blooded	Number legs	Reptile
Cobra	1	1	1	1	0	1
Rattlesnake	1	1	1	1	0	1
Boa constrictor	0	1	0	1	0	1
Chicken	1	1	0	1	2	0
Guppy	0	1	0	0	0	0
Dart frog	1	0	1	0	4	0
Zebra	0	0	0	0	4	0
Python	1	1	0	1	0	1
Alligator	1	1	0	1	4	1

Distance Matrix

	cobra	rattlesnake	boa constrictor	chicken	guppy	dart frog	zebra	python	alligator
cobra	--	0.0	1.414	2.236	1.732	4.243	4.472	1.0	4.123
rattlesnake	0.0	--	1.414	2.236	1.732	4.243	4.472	1.0	4.123
boa constrictor	1.414	1.414	--	2.236	1.0	4.472	4.243	1.0	4.123
chicken	2.236	2.236	2.236	--	2.449	2.646	2.646	2.0	2.0
guppy	1.732	1.732	1.0	2.449	--	4.359	4.123	1.414	4.243
dart frog	4.243	4.243	4.472	2.646	4.359	--	1.414	4.359	1.732
zebra	4.472	4.472	4.243	2.646	4.123	1.414	--	4.359	1.732
python	1.0	1.0	1.0	2.0	1.414	4.359	4.359	--	4.0
alligator	4.123	4.123	4.123	2.0	4.243	1.732	1.732	4.0	--

Producing the Distance Matrix


```
columnLabels = []
for a in animals:
    columnLabels.append(a.getName())
rowLabels = columnLabels[:]
tableVals = []
#Get distances between pairs of animals
#For each row
for a1 in animals:
    row = []
    #For each column
    for a2 in animals:
        if a1 == a2:
            row.append('--')
        else:
            distance = a1.distance(a2)
            row.append(str(round(distance, precision)))
    tableVals.append(row)
```

Producing the Table


```
#Produce table
table = pylab.table(rowLabels = rowLabels,
                    colLabels = columnLabels,
                    cellText = tableVals,
                    cellLoc = 'center',
                    loc = 'center',
                    colWidths = [0.138]*len(animals))
table.auto_set_font_size(False)
table.set_fontsize(10)
table.scale(1, 2.5)
pylab.axis('off')
pylab.savefig('distances')
```


Using Distance Matrix for Classification

- Simplest approach is probably nearest neighbor
- Remember training data
- When predicting the label of a new example
 - Find the nearest example in the training data
 - Predict the label associated with that example



K-nearest Neighbors



Advantages and Disadvantages of KNN

■ Advantages

- Learning fast, no explicit training
- No theory required
- Easy to explain method and results

■ Disadvantages

- Memory intensive and predictions can take a long time
 - Are better algorithms than brute force
- No model to shed light on process that generated data

■ Another method, logistic regression, covered in textbook

An Example of KNN

Features						Label
Name	Egg-laying	Scales	Poisonous	Cold-blooded	Number legs	Reptile
Cobra	1	1	1	1	0	1
Rattlesnake	1	1	1	1	0	1
Boa constrictor	0	1	0	1	0	1
Chicken	1	1	0	1	2	0
Guppy	0	1	0	0	0	0
Dart frog	1	0	1	0	4	0
Zebra	0	0	0	0	4	0
Python	1	1	0	1	0	1
Alligator	1	1	0	1	4	1

Handwritten annotations:

- A bracket above the first five columns (Egg-laying to Number legs) is labeled "Features".
- A bracket above the last column (Reptile) is labeled "Label".
- A red bracket around the "Number legs" column is labeled "dis-proportionate places".
- A red arrow points from the "Number legs" column to the "Reptile" column.
- A red "1/0" is written next to the "Number legs" column.

KNN with $k = 3$

	cobra	rattlesnake	boa constrictor	chicken	guppy	dart frog	zebra	python	alligator
cobra	--	0.0	1.414	2.236	1.732	4.243	4.472	1.0	4.123
rattlesnake	0.0	--	1.414	2.236	1.732	4.243	4.472	1.0	4.123
boa constrictor	1.414	1.414	--	2.236	1.0	4.472	4.243	1.0	4.123
chicken	2.236	2.236	2.236	--	2.449	2.646	2.646	2.0	2.0
guppy	1.732	1.732	1.0	2.449	--	4.359	4.123	1.414	4.243
dart frog	4.243	4.243	4.472	2.646	4.359	--	1.414	4.359	1.732
zebra	4.472	4.472	4.243	2.646	4.123	1.414	<div> <div>→ non reptile</div> <div>reptile</div> </div>		
python	1.0	1.0	1.0	2.0	1.414	4.359			
alligator	4.123	4.123	4.123	2.0	4.243	1.732			

not reptile

An Example of KNN

Features						Label
Name	Egg-laying	Scales	Poisonous	Cold-blooded	Number legs	Reptile
Cobra	1	1	1	1	0	1
Rattlesnake	1	1	1	1	0	1
Boa constrictor	0	1	0	1	0	1
Chicken	1	1	0	1	2	0
Guppy	0	1	0	0	0	0
Dart frog	1	0	1	0	4	0
Zebra	0	0	0	0	4	0
Python	1	1	0	1	0	1
Alligator	1	1	0	1	4	1

KNN with $k = 3$

binary representation of leg

	cobra	rattlesnake	boa constrictor	chicken	guppy	dart frog	zebra	python	alligator
cobra	--	0.0	1.414	2.236	1.732	1.732	2.236	1.0	1.414
rattlesnake	0.0	--	1.414	2.236	1.732	1.732	2.236	1.0	1.414
boa constrictor	1.414	1.414	--	2.236	1.0	2.236	1.732	1.0	1.414
chicken	2.236	2.236	2.236	--	2.449	2.0	2.0	2.0	1.0
guppy	1.732	1.732	1.0	2.449	--	2.0	1.414	1.414	1.732
dart frog	1.732	1.732	2.236	2.0	2.0	--	1.414	2.0	1.732
zebra	2.236	2.236	1.732	2.0	1.414	1.414			
python	1.0	1.0	1.0	2.0	1.414	2.0			
alligator	1.414	1.414	1.414	1.0	1.732	1.732			

reptile ✓

A More General Approach: Scaling

■ Z-scaling

- Each feature has a mean of 0 & a standard deviation of 1

■ Interpolation

- Map minimum value to 0, maximum value to 1, and linearly interpolate

```
def zScaleFeatures(vals):  
    """Assumes vals is a sequence of floats"""  
    result = pylab.array(vals)  
    mean = float(sum(result))/len(result)  
    result = result - mean  
    return result/stdDev(result)  
  
def iScaleFeatures(vals):  
    """Assumes vals is a sequence of floats"""  
    minVal, maxVal = min(vals), max(vals)  
    fit = pylab.polyfit([minVal, maxVal], [0, 1], 1)  
    return pylab.polyval(fit, vals)
```